

Msc Applied Artificial Intelligence**Course: Data-Science-and-Data-Analytics**

ITESM

Data Analysis using Pandas and Python

- Name: Sebastian Ezequiel Coronado Rivera
- Student ID: A01212824

In this activity, you'll use the data file `LaqnData.csv`. Each row in this dataset shows an hourly measurement record of one of the following five air pollutants:

- NO (nitric oxide)
- NO2 (nitrogen dioxide)
- NOX (nitrogen oxides)
- PM10 (suspended particles with an aerodynamic diameter equal to or less than 10 micrometers)
- PM2.5 (suspended particles with an aerodynamic diameter equal to or less than 2.5 micrometers)

Data was collected at one location in London throughout 2017.

IMPORTANT NOTE: Please be sure to answer all questions *explicitly*.

1. Download the file: `LaqnData.csv` and save all its records in a data frame (`air_df`).
- Observe the structure and content of the data frame using the attributes and methods studied (`shape`, `columns`, `head()`, `tail()`, `dtypes`).
- Calculate the percentage of missing values per column.

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import os
#DIR = "/content/drive/MyDrive/Colab Notebooks"
#os.chdir(DIR)

air_df = pd.read_csv('LaqnData.csv')
#air_df.describe()
print('Structure and content of the data frame')
print('\n\n1. Shape \n\n', air_df.shape)
print('\n\n2. COLUMNS \n\n ', air_df.columns)
print('\n\n3. HEAD \n\n', air_df.head())
print('\n\n4. SHAPE \n\n', air_df.tail())
print('\n\n5. DTYPES \n\n', air_df.dtypes)

print('\n\nPercentage of missing values\n\n')
missing_data = air_df.isnull().sum() / len(air_df) * 100
print(missing_data)
```

 Structure and content of the data frame

1. Shape

(43800, 6)

2. COLUMNS

```
Index(['Site', 'Species', 'ReadingDateTime', 'Value', 'Units',
       'Provisional or Ratified'],
      dtype='object')
```

3. HEAD

```
Site Species ReadingDateTime Value Units Provisional or Ratified
```

0	CT3	NO	01/01/2017 00:00	3.5	ug m-3		R
1	CT3	NO	01/01/2017 01:00	3.6	ug m-3		R
2	CT3	NO	01/01/2017 02:00	2.2	ug m-3		R
3	CT3	NO	01/01/2017 03:00	2.1	ug m-3		R
4	CT3	NO	01/01/2017 04:00	3.3	ug m-3		R

4. SHAPE

	Site	Species	ReadingDateTime	Value	Units	Provisional or Ratified
43795	CT3	PM2.5	31/12/2017 19:00	-2.0	ug m-3	R
43796	CT3	PM2.5	31/12/2017 20:00	6.0	ug m-3	R
43797	CT3	PM2.5	31/12/2017 21:00	5.0	ug m-3	R
43798	CT3	PM2.5	31/12/2017 22:00	5.0	ug m-3	R
43799	CT3	PM2.5	31/12/2017 23:00	6.0	ug m-3	R

5. DTYPES

```

Site                object
Species             object
ReadingDateTime     object
Value               float64
Units               object
Provisional or Ratified object
dtype: object

```

Percentage of missing values

```

Site                0.000000
Species             0.000000
ReadingDateTime     0.000000
Value               7.374429
Units               0.000000
Provisional or Ratified 0.000000
dtype: float64

```

2. Get the number of unique values per column (`nunique()`) to answer:

- How many measurements were made?
- How many air pollutants (`Species`) were analyzed?
- Which columns add no informative value because the value is the same throughout the dataset? Remove them from the dataframe.
- The `Units` column also has no informative value. Check it with the `unique()` function and remove it from the dataframe.

```

print('Unique values per column\n')
print(air_df.nunique())

```

```

print('\n\nTotal measurements\n') #Since there is no column that counts the measurements, we will use the row count to relate it to this dat
num_measurements = len(air_df)
print("Approximate number of measurements:", num_measurements)

```

```

print('\n\nContaminants Analyzed:\n')
num_contaminants = air_df['Species'].nunique()
print(num_contaminants)

```

```

print('\n\nColumns with unique value:\n')
columns_uniquevalue = air_df.columns[air_df.nunique() == 1]
print(columns_uniquevalue)
air_df = air_df.drop(columns_uniquevalue, axis=1)

```

```

print('\n\nVerifying units column:\n')
print(air_df['Units'].unique())
air_df = air_df.drop('Units', axis=1)

```

```

print('\n\nVerifying Data frame after cleaning:\n')
air_df.columns

```

Unique values per column

```

Site                1
Species             5
ReadingDateTime     8760
Value               1847
Units               3
Provisional or Ratified 1
dtype: int64

```

Total measurements

Approximate number of measurements: 43800

Contaminants Analyzed:

5

Columns with unique value:

Index(['Site', 'Provisional or Ratified'], dtype='object')

Verifying units column:

['ug m-3' 'ug m-3 as NO2' 'ug/m3']

Verifying Data frame after cleaning:

Index(['Species', 'ReadingDateTime', 'Value'], dtype='object')

3. Count the values by category (`value_counts()`) for the `Species` column and determine if the readings for each contaminant match the total number of measurements.

```
print('Total measurements:\n')
total_measurements = len(air_df)
print(total_measurements)

print('\n\nMissing measurements:\n')
valuenan = air_df[air_df['Value'].isnull()]
countvaluenan = len(valuenan)
print(countvaluenan)

print('\n\nCount removing missing measurements:\n')
nancleaning = air_df.dropna(subset=['Value'])
countafternancleaning = len(nancleaning)
print(countafternancleaning)

print('\n\nCount measurements by contaminant:\n')
pollutant_count = air_df['Species'].value_counts()
print(pollutant_count)

print('\n\nSum of contaminant counts:\n')
count_sum = pollutant_count.sum()
print(count_sum)

if countafternancleaning == count_sum:
    print("\n\nThe contaminant count matches the total number of measurements with no missing values in 'Value'.")
else:
    print("\n\nThere is a difference count of", countvaluenan, "between the contaminant count", count_sum, "and the total measurements removing missing values")

air_df = air_df.dropna(subset=['Value'])
print('\n\nMeasurement count by contaminants after cleaning with NaN values:\n')
clean_contaminant_count = air_df['Species'].value_counts()
print(clean_contaminant_count)

print('\n\nSum of contaminant counts after cleaning:\n')
clean_count_sum = clean_contaminant_count.sum()
print(clean_count_sum)
```

↩ Total measurements:

40570

Missing measurements:

0

Count removing missing measurements:

40570

Count measurements by contaminant:

```
Species
NO      8660
NO2     8660
NOX     8660
PM10    8657
PM2.5   5933
Name: count, dtype: int64
```

Sum of contaminant counts:

40570

The contaminant count matches the total number of measurements with no missing values in 'Value'.

Measurement count by contaminants after cleaning with NaN values:

```
Species
NO      8660
NO2     8660
NOX     8660
PM10    8657
PM2.5   5933
Name: count, dtype: int64
```

Sum of contaminant counts after cleaning:

40570

4. Use the `groupby()` function to determine the average value per contaminant.

```
average_per_pollutant = air_df.groupby('Species')['Value'].mean()
```

```
print("\nAverage value per pollutant:\n\n")
print(average_per_pollutant)
```



Average value per pollutant:

```
Species
NO      15.045115
NO2     38.010185
NOX     61.078661
PM10    22.551704
PM2.5   14.999831
Name: Value, dtype: float64
```

5. The dataset is in a long format. Apply the appropriate function to convert it to a wide format (`ReadingDateTime` as the index and each contaminant in a separate column). Name the resulting data frame `pvt_df`.

```
print("\nConversion to wide format:\n\n")
```

```
pvt_df = air_df.pivot_table(index='ReadingDateTime', columns='Species', values='Value')
```

```
print(pvt_df)
```



Conversion to wide format:

```
Species      NO  NO2  NOX  PM10  PM2.5
ReadingDateTime
01/01/2017 00:00  3.5  30.8  36.2  35.7   NaN
01/01/2017 01:00  3.6  31.5  37.0  28.5   NaN
01/01/2017 02:00  2.2  27.3  30.7  22.7   NaN
01/01/2017 03:00  2.1  23.5  26.8  20.5   NaN
01/01/2017 04:00  3.3  28.0  33.0  22.1   NaN
...          ...  ...  ...  ...  ...
31/12/2017 19:00  0.7  17.5  18.5  16.3  -2.0
31/12/2017 20:00  0.7  17.5  18.6  14.5   6.0
31/12/2017 21:00  0.7  14.1  15.1   8.6   5.0
```

```
31/12/2017 22:00 1.1 22.0 23.6 12.5 5.0
31/12/2017 23:00 0.9 19.4 20.7 10.4 6.0
```

```
[8737 rows x 5 columns]
```

6. Using the `describe()` function, answer:

- What is the highest `NO2` value recorded? What day does it belong to?
- What is the lowest `PM10` value recorded? What day does it belong to?
- What is the median `NO` value? How is it interpreted?
- What is the first quartile of `PM2.5`? What does it mean?

```
print("Statistical description:\n")
pvt_df_description = pvt_df.describe()
print(pvt_df_description)

print("\n\nHighest NO2 value:\n")
max_no2 = pvt_df['NO2'].max()
max_no2_day = pvt_df['NO2'].idxmax()
print(f"Maximum NO2 value: {max_no2}")
print(f"Day of measurement: {max_no2_day}")

print("\n\nMinimum PM10 value:\n")
min_pm10 = pvt_df['PM10'].min()
min_pm10_day = pvt_df['PM10'].idxmin()
print(f"Minimum PM10 value: {min_pm10}")
print(f"Measurement day: {min_pm10_day}")

print("\n\nMedian NO:\n")
median_no = pvt_df['NO'].median()
print(f"Median NO: {median_no}")
print(f"{median_no} represents the middle value of the measured NO data.")

print("\n\nFirst quartile PM2.5\n")
first_quartile_pm25 = pvt_df['PM2.5'].quantile(0.25)
print(f"The first quartile of PM2.5 is: {first_quartile_pm25}")
print(f"{first_quartile_pm25} represents the value below which 25% of the PM2.5 data falls.")
```

↩ Statistical description:

Species	NO	NO2	NOX	PM10	PM2.5
count	8660.000000	8660.000000	8660.000000	8657.000000	5933.000000
mean	15.045115	38.010185	61.078661	22.551704	14.999831
std	26.678565	18.580841	54.584805	15.344755	13.558588
min	-2.000000	2.800000	1.000000	-5.600000	-8.000000
25%	3.100000	23.800000	30.300000	13.900000	7.000000
50%	7.100000	35.300000	46.900000	19.500000	11.000000
75%	15.500000	49.900000	73.800000	27.300000	18.000000
max	401.799990	120.200000	734.299990	633.099980	128.000000

Highest `NO2` value:

Maximum `NO2` value: 120.2
Day of measurement: 24/01/2017 19:00

Minimum `PM10` value:

Minimum `PM10` value: -5.6
Measurement day: 12/09/2017 23:00

Median `NO`:

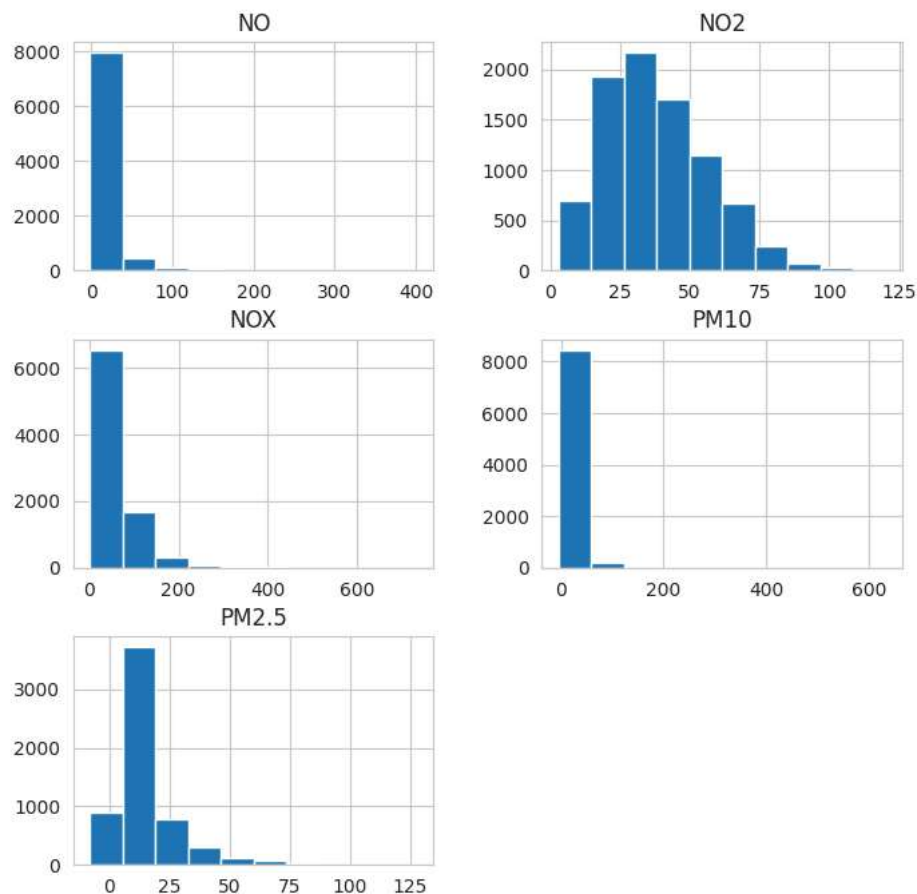
Median `NO`: 7.1
7.1 represents the middle value of the measured `NO` data.

First quartile `PM2.5`

The first quartile of `PM2.5` is: 7.0
7.0 represents the value below which 25% of the `PM2.5` data falls.

7. Draw a histogram for each column of `pvt_df`. Which pollutant has the greatest variability?

```
sns.set_style("whitegrid")
pvt_df.hist(figsize=(8,8))
plt.show()
```



8. Run the following code and comment on what each line does:

```
datetime_df = air_df.ReadingDateTime.str.split(' ',expand=True) #generates a new data frame in which the ReadingDateTime column is split and
datetime_df.columns = ['Date','Time'] # Assign date and time names to the columns of the new data frame generated in the previous line (date
datetime_df # Displays/prints datetime_df
date_df = datetime_df.Date.str.split('/',expand=True) #Similar to the first line, this line splits the date column within datetime_df by the
date_df.columns = ['Day','Month','Year'] #Assigns column names to the date_df data frame (day,month,year)
date_df # Displays/prints date_df
air_df = air_df.join(date_df).join(datetime_df.Time).drop(columns=['ReadingDateTime','Year']) #Joins and deletes columns, air_df.join(date_d
air_df = air_df.set_index(['Month','Day','Time','Species']) # Sets an index of multiple columns defined by month, day time and species, whic
air_df # Displays/prints air_df
```



Value

9. Execute the following statement and compare its output with that of the previous code (exercise 8). Are they the same?

```
air_df.unstack() # Although their results are similar, they are not the same. The data is presented differently. In Exercise 8, four columns
```



Month	Day	Time	Value				
			Species	NO	NO2	NOX	PM10 PM2.5
01	01	01:00	NO	3.6			
		02:00	Value NO	2.2			
		03:00	Species NO	NO	NO2	NOX	PM10 PM2.5
		04:00	NO	3.3			
		00:00		3.5	30.8	36.2	35.7 NaN
		01:00		3.6	31.5	37.0	28.5 NaN
		19:00	PM2.5	-2.0			
		02:00		2.2	27.3	30.7	22.7 NaN
		03:00		2.1	23.5	26.8	20.5 NaN
		21:00	PM2.5	5.0			
12	31	04:00		3.3	28.0	33.0	22.1 NaN
		23:00	PM2.5	6.0			
		19:00		0.7	17.5	18.5	16.3 -2.0
		20:00		0.7	17.5	18.6	14.5 6.0
		21:00		0.7	14.1	15.1	8.6 5.0
		22:00		1.1	22.0	23.6	12.5 5.0
		23:00		0.9	19.4	20.7	10.4 6.0

8737 rows × 5 columns

10. Explain the differences and similarities between the `melt()/pivot()` pair and the `stack()/unstack()` pair. If you had to choose a counterpart to `melt()` between `stack()/unstack()`, which would you choose? Why?

Melt vs. Pivot Similarities Both are used to reshape a DataFrame, either from width to length or vice versa. Both functions offer great flexibility to customize the data transformation through various parameters. Differences Melt is used to transform a DataFrame from a wide format (where each variable is in a different column) to a long format (where there is one column for variables and another for values). It is useful for preparing data for more detailed analysis. Pivot is used to transform a DataFrame from a long format to a wide format. It is useful for creating pivot tables or summarizing data by category.

Stack vs. Unstack

Similarities Both functions are used to reshape a DataFrame, moving data between rows and columns. Both work primarily with DataFrames that have multiple indexes, i.e., multiple levels in the row index.

Differences They are inverse operations. Applying `stack()` and then `unstack()` (or vice versa) to a DataFrame generally returns the original DataFrame. That is, the affected level is different, and so is the direction of movement.

A counterpart to melt?

Well, if it's one of the two, perhaps `unstack` is necessary, since its essence is to return the DataFrame to its original state, but I really think it would be a combination of both.